

29/10

MEMORY SYSTEM

→ DRAM

- For modern memory, one memory access takes 150-200 cycles (not 1 as before)
 - Not counting multiple memory accesses (just one, to fetch it), each instruction will take ~200 cycles
 - Pipelining makes no sense, as a fetched instruction will take only #stages cycles more to complete.

- One possible solution:- 'Cache'

Between processor and memory, add a very small, but fast block of memory. This has, say, only 8 bytes (worth 2 instructions) and is made from more expensive SRAM.

SRAM memory is accessible in single cycle (must fetch consecutively written instructions)

- Fetch 8 bytes (2 instructions) from memory into SRAM (~200 cycles)

Processor fetches	I_1	(200 th cycle)	
	I_2	(201 st cycle)	--- pipelining
	I_3	(400 th cycle)	
	I_4	(401 st cycle)	--- pipelining
- Performance almost doubles by introducing SRAM, as at least 2 instructions are being pipelined.

eg Eg Program for i in range (n):

Assume instruction	I_1
do not access memory	I_2
	I_3
	I_4

If SRAM is 16 bytes, 4 instructions will be fetched in 200 cycles and stored in SRAM.

For every subsequent iteration, we don't need to fetch again and instructions can be pipelined as before

CPI \rightarrow 1



as we are using only 4 instruction multiple times

eg for i in range (100):

if $i \% 2 == 0$:-

$I_1 \ I_2 \ I_3 \ I_4$

else

$I_5 \ I_6 \ I_7 \ I_8$

- With 16 bytes of SRAM, we will have to fetch and replace the four instructions again and again.

$$\text{Total no. of cycles} = (200 + 1 + 1 + 1) \times 100 \quad \text{☹}$$

- With 32 bytes of SRAM, we can get $\text{CPI} \rightarrow 1$, but only if all 8 instructions are consecutively written. Otherwise same as before.

$$(200 + 1 + 1 + 1 + \cancel{1 + 1 + 1 + 1}) \times 50 + (1 + 1 + 1 + 1 + \cancel{1 + 1 + 1 + 1}) \times 49$$

- With 2 SRAM chunks of 16 bytes each, we can get $\text{CPI} \rightarrow 1$. The instructions in each individual chunk must be consecutive.

$$(200 + 1 + 1 + 1 + 200 + 1 + 1 + 1) + (1 + 1 + 1 + 1 + \cancel{1 + 1 + 1 + 1}) \times 48$$

• The SRAM chunks are hidden from programmer and known as 'Cache'.

10/10

→ Bandwidth of Memory :-

$$\frac{\text{Data}}{\text{sec}} = \frac{\text{Data}}{\text{Instruction}} \times \frac{\text{Instruction}}{\text{Cycle}} \times \frac{\text{Cycle}}{\text{Sec}}$$

$$= \left(\frac{\# \text{ Bytes}}{\text{Instruction}} + \frac{\# \text{ Bytes of data}}{\text{Instruction}} \right) \times \text{IPC} \times \text{freq}$$

$$= (4B + 4B) \times 1 \times 3 \times 10^9$$

$$\approx 24 \text{ GBps}$$

- Final CPI depends on no. of 'stalls' required by the program.
- Instructions may stall, for reading from or writing to memory, for 200 cycles
 - Stalling also happens to fetch instructions not present in cache.

- Instructions can be 'hit' or 'miss'
 - present in cache
No penalty
 - not present
200 cycle penalty

• No. of stalls = No. of misses × penalty

$$= \# \text{ Instruction} \times \frac{\# \text{ miss}}{\# \text{ Instruction}} \times \text{penalty}$$

$$= \# \text{ Inst} \times \frac{\# \text{ Memory references}}{\# \text{ Inst}} \times \frac{\# \text{ miss}}{\# \text{ Mem ref}} \times \text{penalty}$$

$$= \text{Inst Count (IC)} \times \frac{\# \text{ Mem ref}}{\# \text{ Inst}} \times \underbrace{\text{Miss rate}}_{\text{In the hands of memory designer}} \times \text{penalty}$$

→ Miss Rate reduction

1 Increase no. of blocks in SRAM

2 Replacing contents of cache :-

Preferentially, do not remove instructions that are consecutively likely to follow current instruction (spatial likelihood) or that are likely to be executed again at some later point in time (temporal likelihood)

- Increasing the number of blocks in SRAM beyond a certain point will increase access time from 1 cycle to 2 cycles. ☹

→ Questions to be answered

1. In which blocks of cache to place?
2. How to identify where required instruction is placed in cache?
3. Which ~~data~~ cache instruction to replace when overwriting?
4. When to write ~~to cache~~ data (obtained from processor) from cache to memory?



In memory, every byte can be addressed as

Block No.	Offset
10 bits (mem)	5 bits
3 bits (cache)	

- We will only be transferring blocks as a whole from memory to cache. After transferring, offset part of each byte remains same. Block no. will be changed accordingly.
- We maintain a mapping of which block no. of memory is ~~map~~ transferred to which block no. of cache.
- Simplest Mapping = 'Direct Mapping'

$$(\text{Mem Block No.}) \% 8 = (\text{Cache Block No.})$$

eg - Mem Block 45 is transferred to Cache block 5.

P.T.O.

i.e. - ~~Tag~~ Mem block

(Cache block no.) = Last 3 bits of (Mem block no.)

\therefore (Cache address) = Last 8 bits of (Mem address)

- The first seven bits of Mem address (which were lost in transferring to cache) will be stored in an adjacent array to cache. (Tag)
- TAG is an 8-element array of 7-bit elements each (one element for each block)

T